

Operating Systems

Courtesy BGU-CSE and Dr. Itamar Cohen

Tutorial 1 – System Calls



*Ben-Gurion University of the Negev
Communication Systems Engineering Department*

Outline

- **Motivation & basics**
- Process control
- Concluding example

Motivation

- A process is ***not supposed*** to access the kernel.
 - It can't access the kernel memory or functions.
- This is strictly enforced (“protected mode”) for good reasons:
 - Can jeopardize other processes running.
 - Cause physical damage to devices.
 - Alter system behavior.
- The system call mechanism provides a safe mechanism to ***request*** specific kernel operations.

System Call - Definition

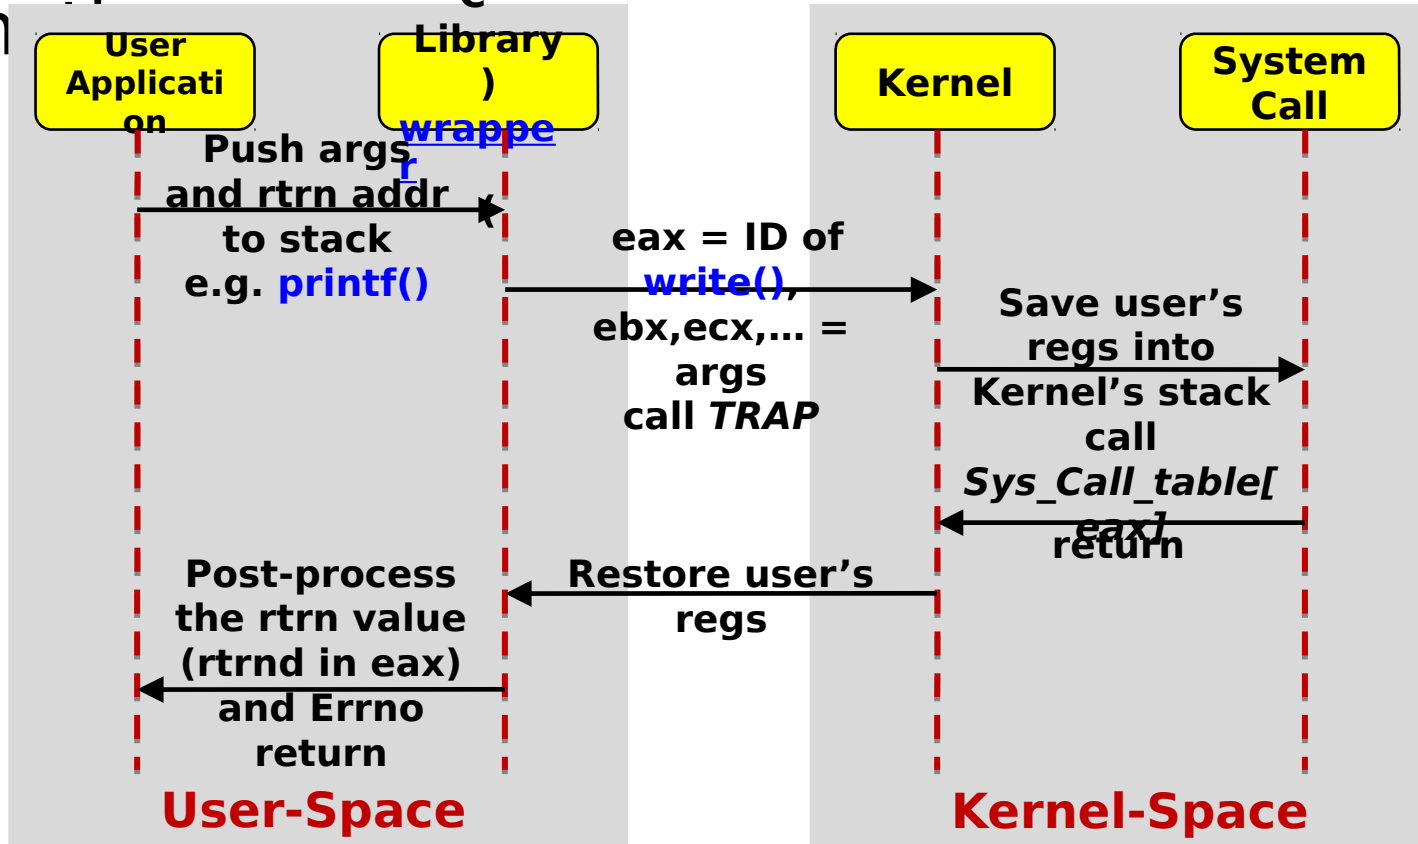
- What is a System Call?
 - An **interface** between a ***user application*** and a ***service*** provided by the operating system (or kernel).
- System call interface – see next slide
 - [More info](#)

System Calls - Categories

- System calls can be roughly grouped into **five** major categories:
 - Process control (e.g. create/terminate process)
 - File management (e.g. read, write)
 - Device management (e.g. logically attach a device)
 - Information maintenance (e.g. set time or date)
 - Communications (e.g. send messages)

System Calls - Interface

- Calls are usually made with C/C++ library functions



Outline

- Motivation & basics
- **Process control**
 - **Creating a new process**
 - Waiting for a process
 - Running a script / command
 - Error report
- Concluding example

Process Control: fork ()

- **pid_t fork(void);**
 - Creates a new process, which is an exact duplicate of the caller
 - Including all file descriptors, registers, instruction pointer, etc
 - Both child and parent resume from after the fork() command
 - and go on their separate ways
 - fork() returns
 - The child's pid, if called by the parent
 - 0, if called by the child
 - A child process can get his parent pid using [getppid\(\)](#)

fork () and Copy On Write: motivation


- When fork is invoked, the parent's information should be copied to its child
- May be wasteful if the child will not need all this information
 - To avoid such situations, use ***Copy On Write*** (COW).

fork(): Example 1

```
int i = 1;
printf("my process pid is %d\n", getpid());
fork_id = fork();
if (fork_id == -1) {
    perror ("Cannot fork\n"); exit
(EXIT_FAILURE); }
else if (fork_id == 0){
    i=7;
    printf("child pid %d, i=%d\n",getpid(),i);
}
else
    printf("parent pid %d, i=%d\n",getpid(),i);
return 0;
```

Output:

```
my process pid is 1000
child pid 1001, i=7
parent pid 1000, i=1
```



Is this the only possible output?
How can we force the output to be deterministic?

Outline

- Motivation & basics
- **Process control**
 - Creating a new process
 - **Waiting for a process**
 - Running a script / command
 - Error report
- Concluding example

Zombies

- When a process ends, the memory and resources associated with it are **de-allocated**.
- However, the entry for that process is **not** removed from its parent's process table.
 - This allows the parent to **collect** the child's exit status.
- When this data is not collected by the parent the child is called a "**zombie**".
 - Such a leak is usually not worrisome in itself. Actually, in some (rare) situations, a zombie is actually **desired** – e.g., for preventing the creation of another child process with the same PID.
 -



Detecting and collecting zombies


- A Zombie can be **collected** by the parent process with the `wait()` system call.
 - See next slide
- Zombies can be **detected** with `ps -el` (marked with 'Z').

wait(), waitpid(), waitid()

- **wait()** – wait for a change in the status of *any* of the children
 - wait – ie, suspend execution of the calling process
 - status: the process terminated / was stopped / was resumed.
 - Once the status of a process is collected, that process is removed from the process table of the collecting process.
- **waitpid(), waitid()** : A finer control than wait(), e.g.
 - Wait for a specific process
 - Wait for any one from a group of processes.
- [Detailed documentation](#)

Orphans

- When the parent process dies before its child, the child becomes an **orphan** process
- The orphans are “adopted” by the **init** system’s process (PID 1)
 - *aka* **reparenting**
- Sometimes orphans are generated intentionally, in order to serve as *daemon* process, e.g.
 - Printer daemon
 - *sshd* – responsible for accepting *secure shell* connections



Outline

- Motivation & basics
- **Process control**
 - Creating a new process
 - Waiting for a process
 - **Running a script / command**
 - Error report
- Concluding example

Running another file: exec ()

- `int execl(char const *path, char const *argv[]);`
- Variants: `int execve(), execvp(), execl(), ...`
- A family of C-library functions, which replace current process image with a new process image (text, data, stack, etc.).
 - Since no new process is created, PID remains the same.
- `exec()` functions ***do not*** return to the calling process unless an error occurred.
 - in which case -1 is returned and `errno` is set with a special value.

Outline

- Motivation & basics
- **Process control**
 - Creating a new process
 - Waiting for a process
 - Running a script / command
 - **Error report**
- Concluding example

errno

- A system variable, which is set by system calls in the event of an error
- Usually *indicates what went wrong*
 - However, *“a function that succeeds is allowed to change errno”* (Linux' manual)
 - The existence or an error is indicated by the function's return value
 - Usually -1 indicates an error
- Frequently a macro.
 - E.g. EACCES (permission denied), EAGAIN (rsrc temporarily unavailable).
- **errno** is *thread local* and *thread-safe*, meaning that setting it in one thread does not affect its value in any other thread.
- Be wary of mistakes such as:
 - ```
if (call() == -1){
 printf("failed...");
 if (errno == ...)
}
```
- Code defensively! Use **errno**

What's the problem?  
*errno* may have been  
changed by *printf()*

# Process control - example 2


```
int main(int argc, char **argv){
 ...
 while(true){
 type_prompt();
 read_command(command, params);
 pid = fork();
 if (pid<0){ //fork failed
 if (errno == EAGAIN) {
 printf("ERROR can't allocate memory\n");
 continue;
 }
 else {
 ... //handle other possible errors
 }
 }
 if (pid>0) //parent
 wait(&status);
 else //child
 execvp(command, parmas);
 }
}
```

# Outline

- Motivation & basics
- Process control
- **Concluding examples**

# Concluding examples: example 5.1

```
int main(int argc, char **argv)
{
 int i;
 for (i=0; i<10; i++){
 fork();
 printf("Hello\n");
 }
 return 0;
}
```




$$\sum_{i=0}^9 2^{i+1} = 2046$$

How many lines of "Hello" will be printed in the following example?

# Concluding examples: example 5.2

```
int main(int argc, char **argv)
{
 int i;
 for (i=0; i<10; i++)
 fork();
 printf("Hello\n");
 return 0;
}
```



$$2^i = 2^i = 1024$$

How many lines of "Hello" will be printed in the following example?