

Operating Systems

Courtesy of BGU-CSE and Dr.
Itamar Cohen

Tutorial 3: Threads



*Ben-Gurion University of the Negev
Communication Systems Engineering Department*

Outline

- **Motivation & basics**
- User-level threads
- Kernel-level threads
- Coding threads

Motivation

- Suppose editing a large WORD doc. This requires
 - User interaction: eg deleting a line and echoing it on the screen
 - Editing the whole document: eg, after a line was deleted, the rest of document is move 1 line upward
 - Automatic disk backups
- Similar scenarios: Editing a .xls, (un)zipping, ...
- One process
 - Too slow
- Many processes:
 - Problem: processes cannot access the same place (file) in memory
 - Long *context switches*
 - High system cost (entries in the process table)
- □ Solution: thread – a refined division of tasks

Threads - General

- One / multiple thread(s) within a process
- Allows multiple independent executions under the same process
- Possible states:
 - Running
 - Ready
 - Blocked
 - (Terminated)

Threads - Advantages

- ☐ **Share** open files, data structures, global variables, child processes, etc.
- ☐ Peer threads can **communicate** without using system calls.
- ☐ Threads are 10-100 times **faster** to create / terminate / switch than processes.

Threads - Disadvantages

- ❑ **Security & stability:** open files, data structures, global variables, child processes etc are shared
- ❑ **Signaling** a thread affects **all threads** of that process

Threads vs. Processes

Threads	Processes
shared data	unique data
shared code	unique code
shared open I/O	unique open I/O
*shared signal table	*unique signal table
unique stack	unique stack
unique PC	unique PC
unique registers	unique registers
unique state	unique state
light context switch	heavy context switch

* Signal handlers must be shared among all threads of a multithreaded app. However, each thread must have its own mask of pending / blocked signals: [use pthread_mask rather than sigprocmask](#).

Implementation dilemmas

- `fork()`

- Duplicate the calling thread / all threads?
- OS dependent
 - Many UNIX systems implement both types of `fork()` (e.g. Solaris 10).
 - In Linux, only the forking thread is duplicated. However, this causes problems, eg in case where the child holds mutex.

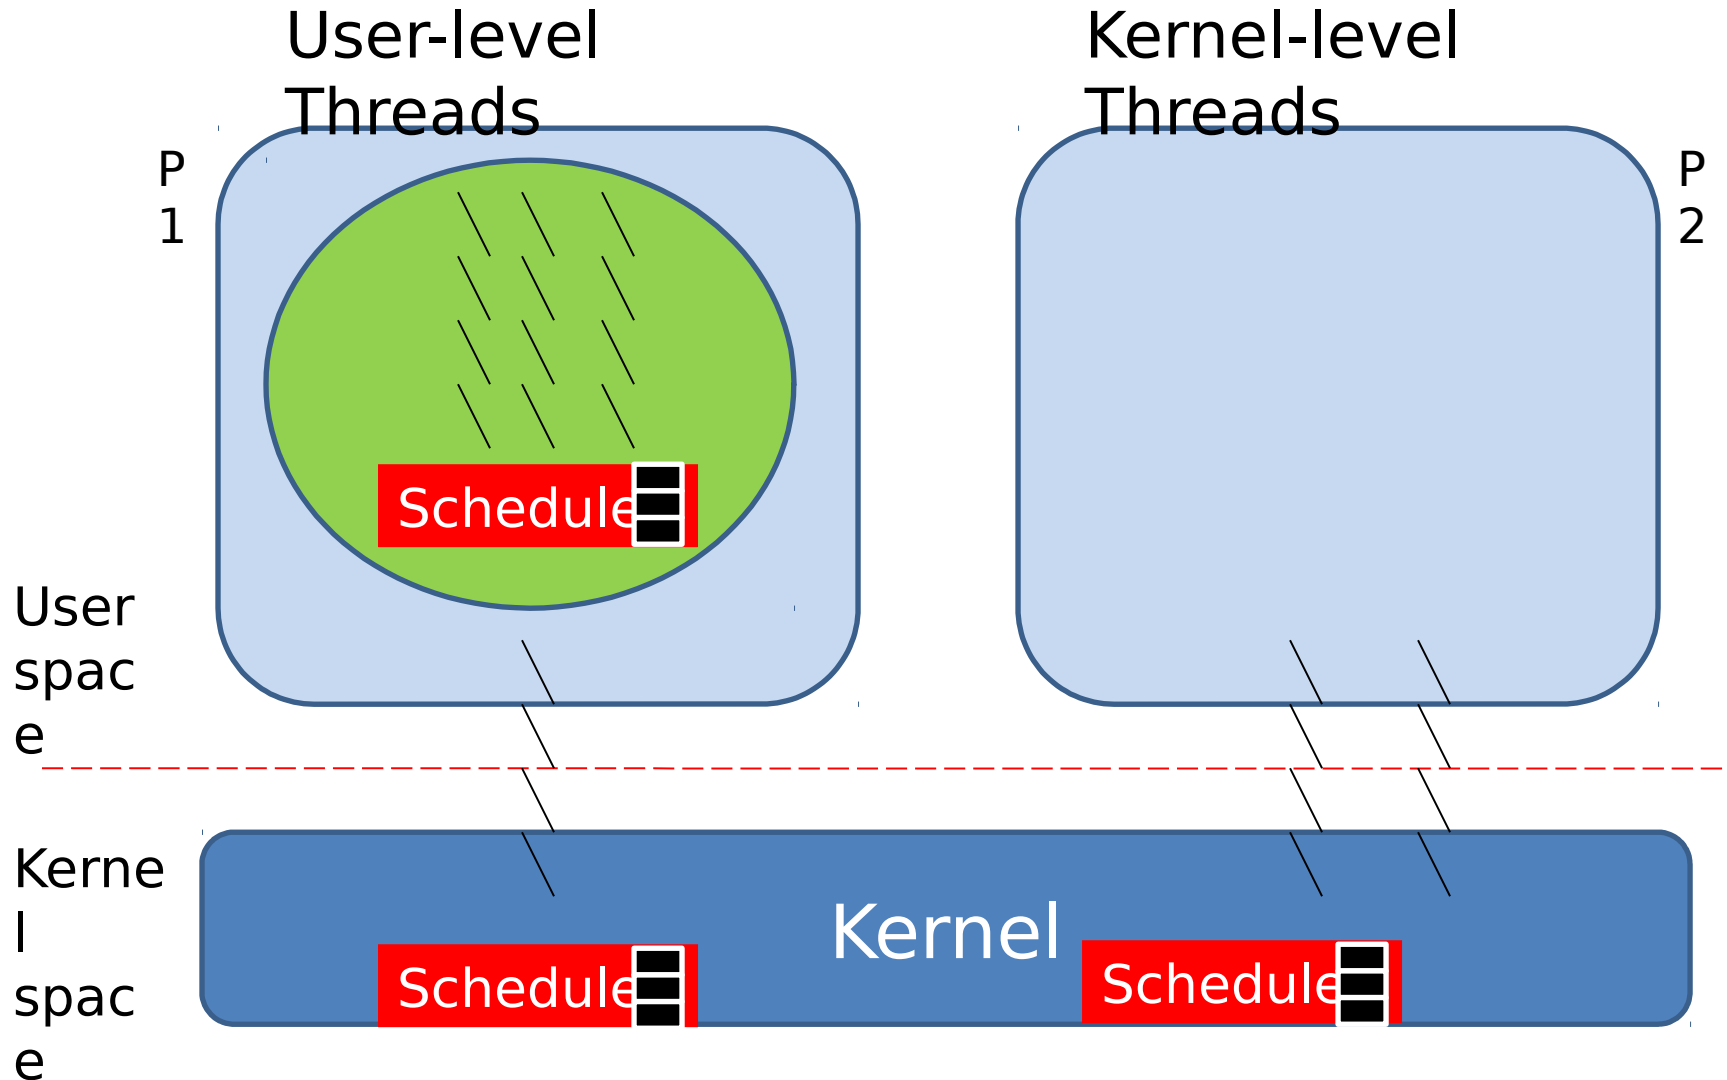
- `exec()`

- Does the command replace the entire process?
 - Yes.

- Divide process into threads by...

- the user / the kernel?

User-Level / Kernel-Level Threads



Outline

- Motivation & basics
- **User-level threads**
- Kernel-level threads
- Coding threads

User-Level Threads

- Implemented in user-level libraries
- Thread switching does not need to call the OS or to cause interrupts.
- The user application schedules the process's CPU time for its internal threads
- Utilizes only a single CPU, as the OS won't allocate multiple CPUs for one process

User-Level Threads - Advantages

- **Compatibility:** Can be implemented on an Operating System that does not support threads.
- **Simple Representation:**
Each thread is represented simply by a PC, registers, stack and a small control block, all stored in the user process address space.
- **Simple & fast management:**
Creating a thread, synchronization and switching between threads can all be done without intervention of the kernel, and are therefore cheaper and ~100 times faster than in kernel-level threads.

User-Level Threads - Disadvantages

- □ Lack of coordination between threads and operating system kernel. Therefore, process as whole gets one time slice irrespective of whether process has one thread or 1000 threads within.
- □ A system call (in one of the thread) causes the OS to block the whole process, even if there are runnable threads left in that processes.
- □ The kernel's inability to notice between user level threads makes it difficult to design preemptive (timed?) scheduling between threads of the same process

Outline

- Motivation & basics
- User-level threads
- **Kernel-level threads**
- Coding threads

Kernel-Level Threads

- All threads are visible to the kernel
- The kernel manages and schedules the threads
- There exist system calls to create and manage threads


Kernel-Level Threads - Advantages

- □ The kernel can smartly schedule between processes with different number of threads
- □ Kernel-level threads are especially good for applications that frequently block
- □ In a multi-processor, a few CPUs can run simultaneously different threads of the same process

Kernel-Level Threads - Disadvantages

- □ Creation, management and switching of threads is MUCH more expensive and slow than user-level threads.

User-Level vs. Kernel-Level Threads

Kernel-level threads	User-level threads	
Visible to the kernel		
Kernel defined		
Preemptive		
Slower, done by the kernel		
a single thread		
kernel	process	Thread table held by...

Outline

- Motivation & basics
- User-level threads
- Kernel-level threads
- **Coding threads**

POSIX Threads & global variables

- Historically, POSIX functions assumed a single thread per process.
 - E.g., consider a naive implementation of `errno` in a multi threaded environment.
- Hence, the need for reentrant functions.
- While this is supported by many standard functions, the compiler must be aware of the need for re-entrant functions:
 - `gcc -D_REENTRANT -lpthread ...`

Threads in POSIX: pthreads

`int pthread_create`

`(pthread_t* thread, pthread_attr_t* attr, void* (*start_func)(void*) , void* arg)`

Creates a new thread of control that executes concurrently with the calling thread.

On success, the identifier of the newly created thread is stored in the location pointed by the *thread* argument, and a 0 is returned. On error, a non-zero error code is returned.

attr enables applying attributes to the new thread (e.g. detached, scheduling-policy). Can be NULL (default attributes).

start_func is a pointer to the function the thread will start executing. The function receives one argument of type *void** and returns a *void**.

arg is the parameter to be given to *func*.

`pthread_t pthread_self ()`

Returns this thread's identifier.

Threads in POSIX (pthreads) – cont.

int **pthread_join** (pthread_t th, void** thread_return)

Suspends the execution of the calling thread until the thread identified by *th* terminates.

On success, the return value of *th* is stored in the location pointed by *thread_return*, and a 0 is returned. On error, a non-zero error code is returned.

At most one thread can wait for the termination of a given thread. Calling **pthread_join** on a thread *th* on which another thread is already waiting for termination returns an error.

th is the identifier of the thread that needs to be waited for

thread_return is a pointer to the returned value of the *th* thread (can

void **pthread_exit** (void* ret_val)

Terminates the execution of the calling thread. Doesn't terminate the whole process if called from the main function.

If *ret_val* is not null, then *ret_val* is saved, and its value is given to the thread which performed *join* on this thread; that is, it will be written to the *thread_return* parameter in the **pthread_join** call.

Example files

- `thread_hello_world.c`
- `thread_exercise2.c`